

000
001
002
003
004
005
006
007
008
009
010
011

Figure 1. Turning that frown upside down: We propose that capsules can model action units as individual part deformations

Abstract

In this paper we motivate the use of capsule networks for facial action unit detection. We argue that action unit activations may be seen as local part deformations - for example AU 1,2, and 4 are deformations of the part 'eyebrows'. Different part deformations in a regular convolutional network must be modeled and represented as separate neurons. However, with capsule networks a part can be represented by a single capsule, and its deformations can be modeled by its direction. We test this hypothesis by creating a capsule network for action unit recognition. We find that capsule networks are indeed able to model action units and local part deformations as well. These lead to state-of-the-art results on the BP4D and DISFA datasets. We analyse the learned capsules' properties and find that capsule magnitude correlates with expression intensity and that capsule pose captures varied attributes such as face size, lighting, pose, and skin color. Finally we use activation gradient ascent to visualize capsule direction, and find that a single capsule can represent multiple deformations of the same part, while a single convolution neuron does not.

1. Introduction

Facial Action Coding System (FACS) is a system to define and name facial movements by their appearance on the face. Informed by the underlying muscular structure of a

face, FACS annotation can be reliably used for describing as well as identifying facial expressions and is therefore not as subjective as grimace scales. There are 24 main facial action units to describe the human face. Additionally, action units can also be coded for intensity on a 4 to 5 point scale. While, extremely useful, manual action unit coding is a cumbersome process that can only be carried out by trained experts. Due to this hurdle, automatic action unit detection is an important problem for computer vision research.

Facial action unit detection requires identifying subtle deformations on parts of the face. Consequently, features that capture local movements around key parts of the face have been used to train machine learning systems. For example, in [1], a seminal work on emotion and facial action unit understanding, Gabor features were extracted around keypoints of the face to capture local muscle deformations. Developing features that capture part deformations well has also motivated more recent work [43, 19] where separate convolutional filters are trained to correspond to different parts of the face. The motivation behind these works is similar – the better we are able to model how parts of a face look and change, the better we can detect action units.

Recently, capsule networks [29, 13] were introduced. There are two primary reasons why capsule architectures have advantages over regular convolutional neural networks. Intuitively, neurons in a CNN can represent an attribute of the input image – such as the presence of an eye – and the activation of a neuron represents a confidence value in whether that attribute can be found in the image. With capsules the expressive power is increased - its activation can represent a confidence value in its presence or absence, and its direction can represent properties of the attribute. For example, the direction of the capsule can indicate how rotated the eye is, whether it is open or close, etc. The second advantage of capsules is that the additional representative capacity allows for complex routing procedures. The pose of a capsule can be used to determine how it is propagated through the network. This is in contrast to neuron activations in a CNN that are propagated solely on the basis of its scalar value. As a result, the routing procedure followed

108 by capsule networks can mimic the effect of a much deeper
109 convolution neural network, trained with various data trans-
110 formed augmentation techniques.
111

112 We believe that the higher expressive power capsules af-
113 ford to each visual attribute of an image can directly trans-
114 late to better modeling of local part deformations. Action
115 unit activations can be seen as local part deformations, and
116 therefore capsules can be better at detecting and modeling
117 action units. To give a naive parallel example, if a capsule
118 learns to detect lips, its pose can represent the type of defor-
119 mation the lips are in – so different capsule poses can end
120 up representing action units 12, 14, 15. At the same time,
121 for a convolutional network, if a neuron comes to be asso-
122 ciated with lips, it cannot express the pose the lips are in. It
123 can simply be either active - indicating the presence of lips
124 – or inactive – indicating that the lips are not present. In
125 order to represent lips in AU-12 or in AU-15, the network
126 must learn to associate separate neurons.
127

128 **Contributions.** Our main contributions in this paper are
129 as follows:
130

131 We present results that indicate that capsules are indeed
132 better than CNNs at modeling local part deformations - and
133 therefore action units. Previous work [29, 13] have shown
134 that capsules can model global deformations - so the net-
135 work can generalize well across image level deformations
136 (affine transformations of MNIST), or viewpoint deformations
137 (azimuth changes on small-NORB). In this paper, we
138 show that capsules can also capture small deformations well
139 - and may therefore be extremely useful in other areas such
140 as fine-grained classification.

141 We develop a capsule network for action unit detection
142 that gives state-of-the-art results across two large action unit
143 datasets. On BP4D dataset we outperform the closest base-
144 line architecture by 14.1% in AUC. We replicate similar
145 performance gains on DISFA. We additionally present re-
146 sults on emotion recognition for CK+ dataset, and find that
147 our architecture generalizes well.

148 To the best of our knowledge, capsule networks have not
149 been used to perform facial action unit detection before. To
150 this end, we thoroughly analyse and visualize the learned
151 capsule networks. We visualize the effect of changing cap-
152 sule magnitude as well as capsule direction via a reconstruc-
153 tion network. We find that capsules are able to model face
154 pose, shape, lighting, and skin color, and that capsule mag-
155 nitude is correlated with action unit or emotion intensity. Fi-
156 nally, we use activation maximizing gradient ascent to visu-
157 alize capsule features and compare them with regular con-
158 volution networks. We find that capsules are able to model
159 part deformations as changes in capsule direction, where as
160 individual convolution neurons are not able to model multi-
161 ple part deformations.

2. Related Work

2.1. Capsule Networks

Capsule networks were first proposed in [29]. The net-
work replaces scalar neurons with higher dimension cap-
sules - so that activation and neuron attributes can be mod-
eled jointly. In addition, capsule direction or pose can be
used to route capsules between higher layers - which re-
places pooling based routing in convolutional networks. In
[13], the authors propose vector capsules whose magnitude
represents the activation of a capsule. Iterative routing is
done using a simple agreement between lower and higher
level capsule directions. In [13] the authors introduce ma-
trix capsules, where a separate value represents the cap-
sule activation. Routing is done using an EM algorithm,
such that the probability distributions of higher and lower
level capsules between consecutive layers are in agreement.
Furthermore, the authors introduce convolutional capsules,
whereas [29] only worked with fully-connected capsules.
In this paper, we use vector capsules, with dynamic routing,
and work with fully-connected capsules only.

2.2. Facial Action Unit Understanding

Papers in action unit understanding have focused on two
broad sub problems - action unit intensity estimation, and
action unit detection.

A number of traditional non-deep approaches improve
action unit understanding by exploiting the co-occurrence
patterns between action units - either by developing a
learning model that can help capture inter-AU relations
[34, 37, 7], by developing a model based on prior knowl-
edge of AU relationships and semantics [32, 20], or by using
a data-driven approach to learn important AU relationships
[46]. In particular, [42] jointly learns to identify important
patches, and positive and negative correlations between ac-
tion units for understanding action units.

Traditional approaches have also learned action units by
assistance from facial keypoints - features extracted around
keypoints are used for action unit detection. Some examples
of such approaches are [7, 18, 3, 33, 36].

Deep learning has also been applied to the problem of ac-
tion unit detection with great success [9, 10, 14, 11, 43, 19].
Amongst these, two papers in particular require discussion.
In [43] the authors develop a ‘region layer’ that splits the
incoming convolution map into a grid and develops sepa-
rate convolutional maps for each grid section. The result-
ing map is concatenated spatially and propagated through
the network. In a similar vein, [19] also explicitly design
their deep network to develop features for parts of a con-
volutional map - however, unlike [43] the spatial regions that
are broken up are based on facial keypoint locations and
their correlations with action units. Both methods explore
a similar idea - to develop separate features for parts of a

162
163
164
165
166
167
168
169
170
171
172
173
174
175
176
177
178
179
180
181
182
183
184
185
186
187
188
189
190
191
192
193
194
195
196
197
198
199
200
201
202
203
204
205
206
207
208
209
210
211
212
213
214
215

216 face - as is based on the intuition that different areas of the
 217 face correspond to different AU activations that require their
 218 own unique set of features for identification.
 219

220 ECCV papers [5, 30]. They both do facial action unit
 221 recognition.

222 CVPR papers [40, 28] do weakly supervised au recogni-
 223 tion. [12] optimizes filter size per au by either expanding
 224 or contracting filter sizes from a base size over training.
 225 the model relies on separate models for each action unit.
 226 [39] uses expression indenpendant and expression depen-
 227 dant prior knowledge about action units to learn au classi-
 228 fiers without direct supervision. [41] does au intensity esti-
 229 mation.

230 2.3. Expressions

231 Facial expression prediction is a well-explored topic of
 232 research in computer vision. We primarily focus on action
 233 unit detection, but also show qualitative and quantitative
 234 results on expression detection. Some approaches that do not
 235 use deep learning are, [1, 8, 31, 45], of which [1] is of partic-
 236 ular note for creating a pipeline based on extracting features
 237 around facial keypoints, detecting action units, and fusing
 238 action unit detections temporally for emotion detection.
 239

240 A number of papers also explore emotion understanding
 241 in a simple deep feed-forward classification network setting
 242 [23, 16, 26, 44, 6]. [21, 22] attempt to enforce AU under-
 243 standing to the end of emotion classification. Of these, [16]
 244 is notable for impressive results on expression detection and
 245 demonstrating the importance of data augmentation for the
 246 task of expression understanding. [6] is also an important
 247 paper that proposes a two-stage training pipeline to transfer
 248 VGG-Face [27] features for the task of expression classifi-
 249 cation. Also noteworthy is [15], which uses facial keypoint
 250 locations over time to train a network that is meant to cap-
 251 ture temporal deformations alongside a traditional image-
 252 based CNN. Lastly, [17] proposes an encoder-decoder type
 253 architecture that learns from pairs of neutral/non-neutral ex-
 254 pressions to develop features that are discriminative for ex-
 255 pression classification.
 256 CVPR [35].

257 3. Approach

258 Our network comprises two modules - a capsule network
 259 that outputs action unit capsules, and a reconstruction net-
 260 work that takes concatenated action unit capsules, and is
 261 trained to reconstruct the input image. During training, the
 262 capsules for all classes apart from the ground-truth classes
 263 are zeroed-out and used as input to the reconstruction net-
 264 work. In this way, the reconstruction network does not, di-
 265 rectly, affect classification accuracy.
 266

267 We train our network with color inputs of size 96×96 .
 268 We train with three routing iterations.
 269

270 The capsule network architecture comprises of two con-
 271 volution layers, with 64 and 128 filters, and kernel size of 5.
 272 Each is followed by max-pooling and ReLu. The convolu-
 273 tion layers are followed by a primary caps layer with 32 cap-
 274 sules of dimension 8, filter size 7, and stride 3. The resulting
 275 activation map is then fully connected to our class capsule
 276 layer with n capsules, each with dimension 32, where n is
 277 the number of output classes.
 278

279 Our reconstruction network comprises of 3 fully con-
 280 nected linear layers of dimension 512, 1024, and 1024. The
 281 last layer is reshaped to 32×32 and then bilinearly upsam-
 282 pled to 96×96 . In experiments we refer to this model as
 283 ‘Ours’.

284 Additionally, we propose a larger capsule network with
 285 VGG convolution layers as its base. The network is iden-
 286 tical to VGG-16 up to the end of its convolution layers.
 287 The last max-pooling layer is removed. This is followed
 288 by 32 primary capsules of size 8, kernel size of 3 and stride
 289 of 2. The class capsules have dimension 32. The recon-
 290 struction network comprises of three fully-connected lay-
 291 ers with dimensions 512, 1024, and 9408. The output is
 292 resized to 56×56 and then bilinearly upsampled to 224 –
 293 the input image size. We show results of this model with
 294 the convolution layers initialized with both Imagenet
 295 pretrained weights (Ours-VGG) and VGG-Face pretrained
 296 weights (Ours-VGGF). We found Ours-VGGF was prone
 297 to overfitting. We therefore add spatial dropout at 70% after
 298 the last convolution layer.

299 We use the margin loss from [29] to train our networks.
 300 The original loss for class c :

$$L_c = T_c \max(0, m^+ - \|v_c\|)^2 + \lambda(1-T_c) \max(0, \|v_c\| - m^-)^2$$

301 where $T_c = 1$ iff class c is present, $m^+ = 0.9$ and $m^- =$
 302 0.1, v_c is the class capsule, and λ is a downweighting term
 303 for negative samples set to 0.5.

304 For single class classification (such as expression clas-
 305 sification) a softmax operation is applied across all $\|v_c\|$.
 306 However, for the multiclass classification setting (action
 307 unit detection), the softmax is not applied. Note that due to
 308 the capsule squashing operation, the magnitude of all output
 309 capsules still lies between 0 and 1.
 310

311 For action unit detection, the occurrence of different AUs
 312 is highly imbalanced. We therefore modify the margin loss
 313 to work with a class specific weight:

$$L_c = w_c(T_c \max(0, m^+ - \|v_c\|)^2 + \lambda(1-T_c) \max(0, \|v_c\| - m^-)^2)$$

315 where w_c is the inverse of the frequency of an action unit
 316 occurrence, normalized across all classes to sum to one.
 317

318 The margin loss is then averaged across all action unit
 319 instances in a batch, and added with average reconstruction
 320 loss for the batch. We use mean square error to supervi-
 321 se the reconstruction network. The final loss term is:
 322

$$L_{final} = L_{cls} + \alpha L_{recon}$$

324

325

326

327

328

329

330

331

332

333

334

335

336

337

338

339

340

341

342

343

344

345

346

347

348

349

350

351

352

353

354

355

356

357

358

359

360

361

362

363

364

365

366

367

368

369

370

371

372

373

374

375

376

377

AU	LSVM[43]	JPMI[42]	DRML[43]	CPM[50]	CNN+LSTM[4]	FVGG[19]	ROI[9]	FERA[14]	Ours	Ours-VGG	Ours-VGGF
1	33.2	32.6	36.8	43.4	31.4	27.8	36.2	28.0	46.8	40.0	47.3
2	22.8	25.6	41.8	40.7	31.1	27.6	31.6	28.0	29.1	27.7	39.9
4	23.1	37.4	43.0	43.4	71.4	18.3	43.4	34.0	52.2	42.2	52.8
6	27.2	42.3	55.0	59.2	63.3	69.7	77.1	70.0	75.3	76.1	77.9
7	47.1	50.5	67.0	61.3	77.1	69.1	73.7	78.0	77.6	71.8	79.9
10	77.2	72.2	65.1	54.1	45.0	78.1	85.0	87.3	82.0	81.8	84.0
12	63.7	74.1	65.8	68.5	82.6	63.2	87.0	78.0	85.0	87.3	88.1
14	64.3	65.7	54.1	52.5	72.9	36.4	62.6	75.0	65.7	63.5	67.2
15	18.4	38.1	36.7	34.0	33.2	26.1	45.7	20.0	33.7	36.1	49.2
17	33.0	40.0	48.0	54.3	53.9	50.7	58.0	36.0	60.4	62.1	65.4
23	19.4	30.4	31.7	39.5	38.6	22.8	38.3	41.0	36.9	35.3	47.7
24	20.1	42.3	30.0	37.8	37.0	35.9	37.4	43.1	44.3	44.6	55.1
Avg	35.3	45.9	48.3	50.0	53.2	43.8	56.4	51.7	57.4	55.7	62.9

Table 1. F1-Frame results on BP4D dataset. Our method outperforms all methods that are trained from scratch and even outperforms FVGG despite not using any external data.

322

AU	LSVM[43]	APL[43]	DRML[43]	FVGG[19]	Ours	Ours-VGG	Ours-VGGF
1	10.8	11.4	17.3	32.5	17.6	15.7	15.7
2	10.0	12.0	17.7	24.3	18.8	25.7	25.7
4	21.8	30.1	37.4	61.0	50.1	41.3	41.3
6	15.7	12.4	29.0	34.2	44.8	52.8	52.8
9	11.5	10.1	10.7	1.67	21.6	40.7	40.7
12	70.4	65.9	37.7	72.1	65.1	70.1	70.1
25	12.0	21.4	38.5	87.3	68.8	62.5	62.5
26	22.1	26.9	20.1	07.1	45.4	47.7	47.7
Avg	21.8	23.8	26.7	40.2	41.5	44.6	44.6

Table 2. F1-Frame results on DISFA dataset. Our method outperforms all methods that follow a similar training protocol as ours.

where α is a weight parameter, L_{recon} is the average reconstruction loss and L_{cls} is the average margin loss. We set α to bring the order of the average reconstruction loss in the same order of magnitude as the averaged margin loss at the beginning of training. We set α to $1e-7$ at for the smaller network and $1e-8$ for the larger VGG based model.

4. Experiments

4.1. Action Unit Detection

4.1.1 Datasets

We present results on two widely-used datasets.

BP4D [38]: The dataset contains 328 videos of 31 subjects while completing eight different tasks designed to elicit emotion. Frames are annotated with 12 different action units. In total there are a little less than 140000 frames that we can use. Following common procedure, we do 3 fold cross validation on subjects, train on 2 folds, and test on the third. Results are collated across folds and reported.

DISFA [25]: 26 subjects are recorded while watching videos. Action units and their intensity are annotated for each frame. Similar to BP4D we conduct 3 fold cross validation, and collate results across folds.

For both datasets we detect and align faces using [2]. Images are randomly horizontally flipped, rotated, scaled, translated, cropped, and pixel augmented for data augmentation.

4.1.2 Metrics

We report F1-Frame score, as well as AUC. The F1 score is the harmonic mean of precision and recall, and used by AU

AU	LSVM	JPMI	DRML	AlexNet	ConvNet	LCN	DRML	Ours	Ours-VGG	Ours-VGGF
1	20.7	40.7	34.9	49.4	51.9	55.7	65.7	60.3	66.1	66.1
2	17.7	42.1	25.8	51.3	50.9	54.5	56.0	54.9	63.7	63.7
4	22.9	46.2	36.1	47.4	53.6	58.8	70.2	62.7	71.8	71.8
6	20.3	40.0	48.3	52.2	53.2	56.6	71.3	74.0	73.7	73.7
7	44.8	50.0	54.3	64.8	63.7	61.0	60.6	58.7	68.5	68.5
10	73.4	75.2	54.3	61.4	62.4	53.6	70.8	73.4	70.1	70.1
12	55.3	60.5	50.0	60.2	61.6	60.8	74.6	81.2	80.8	80.8
14	46.8	53.6	47.7	29.8	58.8	57.0	56.7	56.9	57.5	57.5
15	18.3	50.1	34.9	50.6	49.9	56.2	59.4	71.2	71.2	71.2
17	36.4	42.5	48.5	53.5	48.4	50.0	66.1	67.4	71.0	71.0
23	19.2	51.9	40.5	49.5	50.3	53.9	61.6	60.5	69.6	69.6
24	11.7	53.2	31.7	52.5	47.7	53.9	67.6	68.3	77.4	77.4
Avg	32.2	50.5	42.2	51.8	54.4	56.0	65.0	64.9	70.1	70.1

Table 3. AUC scores on BP4D. We outperform all methods with a large margin.

AU	LSVM	APL	AlexNet	ConvNet	LCN	DRML	Ours	Ours-VGG	Ours-VGGF
1	21.6	32.7	47.8	44.2	44.1	53.3	57.3	56.0	56.0
2	15.8	27.8	52.1	37.3	52.4	53.2	58.9	61.5	61.5
4	17.2	37.9	44.0	47.9	47.7	60.0	70.4	63.2	63.2
6	08.7	13.6	44.3	38.5	39.7	54.9	67.6	70.2	70.2
9	15.0	64.4	48.7	49.5	40.2	51.5	66.0	67.5	67.5
12	93.8	94.2	55.3	54.8	54.7	54.6	77.1	79.0	79.0
25	03.4	50.4	50.2	48.4	48.6	45.6	75.7	72.2	72.2
26	20.1	47.1	45.8	45.8	47.0	45.3	69.2	67.2	67.2
Avg	27.5	46.0	49.1	45.8	46.8	52.3	67.8	67.1	67.1

Table 4. AUC scores on DISFA. We outperform all methods with a large margin.

4.1.3 Implementation Details

For ‘Ours’ on BP4D we train for ten epochs and learning rate of $1e-4$. For finetuning on DISFA, we transfer convolution features only, and train at learning rate of $1e-4$ for the first 5 epochs, and drop it to $1e-5$ for the remaining 5 epochs.

For ‘Ours-VGGF’ we follow the procedure from ROI and fix all convolution layers up to the beginning of conv5. We finetune the conv5 layers at a 10 times lower learning rate, and train for 5 epochs. The capsule, and reconstruction layers are trained at $1e-4$ learning rate. For ‘Ours-VGG’, we additionally finetune the first four convolution layers, and train for 10 epochs. For DISFA finetuning we initialize the convolution layers’ weights with the best performing BP4D model and do XXX. We additionally present results without any BP4D transfer for DISFA (Ours-VGGFS, and Ours-VGGS) that uses the same training procedure as used for BP4D dataset.

4.1.4 Results

Results on BP4D dataset are shown in Table 1 for F1 and Table 3 for AUC. ‘Ours’ shows results after training our model from scratch and is comparable to all columns apart from ROI - which builds on VGG-Face features, and FVGG which is VGG-Face finetuned for AU detection. Our method outperforms all methods that are trained from scratch and even outperforms FVGG despite not using any external data. Overall, our best method outperforms the

432	Method	Accuracy
433	AURF[21]	92.2
434	AUDB[22]	93.7
435	Khorrami[16]	96.4
436	GCNet*[17]	97.28
437	FN2EN*[6]	96.8
438	Ours	93.7
439	Ours-Max	96.2

442 closest baseline on BP4D by 6.5% F1 score with external
 443 data, and our method without any external data outperforms
 444 the closest comparable baseline ([4]) by 4.2% despite not
 445 using any temporal information.

446 Following common procedure we present DISFA results
 447 after transferring from BP4D model in Table 2 for F1 and
 448 Table 4 for AUC. For all models we outperform the base-
 449 lines significantly. We believe we are able to outperform
 450 baselines that often have more parameters because we are
 451 using capsule networks that are better equipped to model
 452 small facial deformations.

4.2. Emotion Detection

455 We also explore the use of capsules on the related task of
 456 emotion recognition on the Cohn-Kanade dataset [24]. We
 457 follow the established protocol of 10 fold cross validation,
 458 and average results across folds.

459 Figure ?? (left) shows our results on 8 emotion classifi-
 460 cation against several state-of-the-art methods. Our results
 461 are comparable to the state-of-the-art. We found that test
 462 results were prone to fluctuate throughout training, and due
 463 to the small dataset size, some folds were prone to overfit.
 464 We therefore also report the best test accuracy we achieve
 465 during training as ‘Ours Max’ to provide an idea of the up-
 466 per limit our model may achieve with more careful training
 467 and hyperparameter searching.

4.3. Visualizing Capsules by Reconstruction

471 Every class capsule is a 32 length long vector. This vec-
 472 tor can be modified by rescaling its magnitude, or altering
 473 its direction. For every altered version of an input image’s
 474 correct class capsule, we can use reconstruction network to
 475 visualize the capsule.

4.3.1 Magnitude

477 Since capsule magnitude represents the confidence of our
 478 network in a capsule class’s presence in the image, we ex-
 479 expect increasing capsule magnitude to create reconstruc-
 480 tions that represent that class even more. In other words,
 481 we would expect the reconstruction of a ‘surprise’ capsule with
 482 less magnitude to show less surprise than a reconstruction
 483 of the same capsule scaled to a higher magnitude.



486
 487
 488
 489
 490
 491
 492
 493
 494
 495
 496
 497
 498
 499
 500
 501
 502
 503
 504
 505
 506
 507
 508
 509
 510
 511
 512
 513
 514
 515
 516
 517
 518
 519
 520
 521
 522
 523
 524
 525
 526
 527
 528
 529
 530
 531
 532
 533
 534
 535
 536
 537
 538
 539

Figure 2. We show the effect of altering magnitude of a class capsule. As the magnitude or activation of a capsule increases, the intensity of the facial expression also increases.



509
 510
 511
 512
 513
 514
 515
 516
 517
 518
 519
 520
 521
 522
 523
 524
 525
 526
 527
 528
 529
 530
 531
 532
 533
 534
 535
 536
 537
 538
 539

Figure 3. Reconstructions of test images with increasing capsule magnitude (left to right) on the BP4D dataset

531
 532
 533
 534
 535
 536
 537
 538
 539

Figure 2 shows reconstructions of different expression capsules with increasing magnitude. From left to right the magnitude was increased from 0.1 to 0.9 at increments of 0.1. The extreme left image is the input image. For each of the capsules, we see the expression become more pronounced and exaggerated. For example, for happiness (row 1), the smile in subsequent reconstructions becomes wider. Similarly, for surprise (row 3), a dark spot resembling an open mouth begins to appear, eventually resembling a full jaw drop.

531
 532
 533
 534
 535
 536
 537
 538
 539

We repeat this experiment with our second network, this time increasing the magnitude of all action units that are present in the image. For BP4D Figure 3(top) increasing magnitude results in unique features of the image being exaggerated, such as skin color (3rd row), face and neck shape (4th and 5th row), or pose (2nd row). At the same time, certain AUs also become prominent. The open mouth smile in the 3rd row, the jaw drop in the 4th row, the furrowed brow in the 2nd row, and the raised chin in the 1st row, all become more obvious as magnitude is increased.

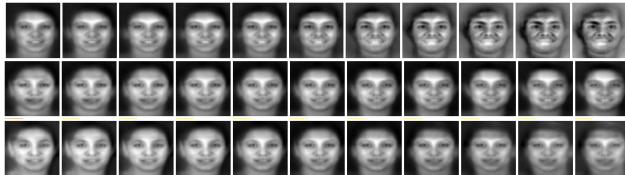
531
 532
 533
 534
 535
 536
 537
 538
 539

However, these results are not as clear as the ones from the expression network, and the changes in action units’ prominence would be hard to see unless we knew what to look for. This can be due to the training data where expression changes are spontaneous and subtle and not posed and exaggerated as in the Cohn-Kanade dataset. They are unlikely to effect the reconstruction loss function enough to create a strong supervision signal. Additionally, there are multiple action units present in every image, so the net-



540
541
542
543
544
545
546
547
548
549
550
551
552
553
554
555
556
557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593

Figure 4. We reconstruct an image for its ground truth class while changing its capsule direction by altering the value of each of its 32 dimensions for the CK+ dataset. The capsule dimensions are associated with attributes as face scale (first row), face shape (second row), and lighting(third row).



557
558
559
560
561
562
563
564
565
566
567
568
569
570
571
572
573
574
575
576
577
578
579
580
581
582
583
584
585
586
587
588
589
590
591
592
593

Figure 5. We reconstruct an image for its ground truth class while changing its capsule direction by altering the value of each of its 32 dimensions for the BP4D dataset. The first row shows face features and skin change. In the second, teeth appear. The neck in the last row indicates that the capsules have also learned pose attribute.

work may only prioritize the reconstruction of obvious action units that create large changes in the face, and not all action units.

4.4. Direction

We can also keep capsule magnitude stable, while changing its direction. For this, we vary the value of each of its 32 dimensions between -0.5 and 0.5 , and reset the capsule magnitude to its original magnitude. In Figures 4 and 5 we show the effect of changing some capsule directions. The capsule dimensions are associated with attributes as varied as face shape, skin color, pose, or the visibility of teeth.

4.5. Visualizing by Gradient Ascent

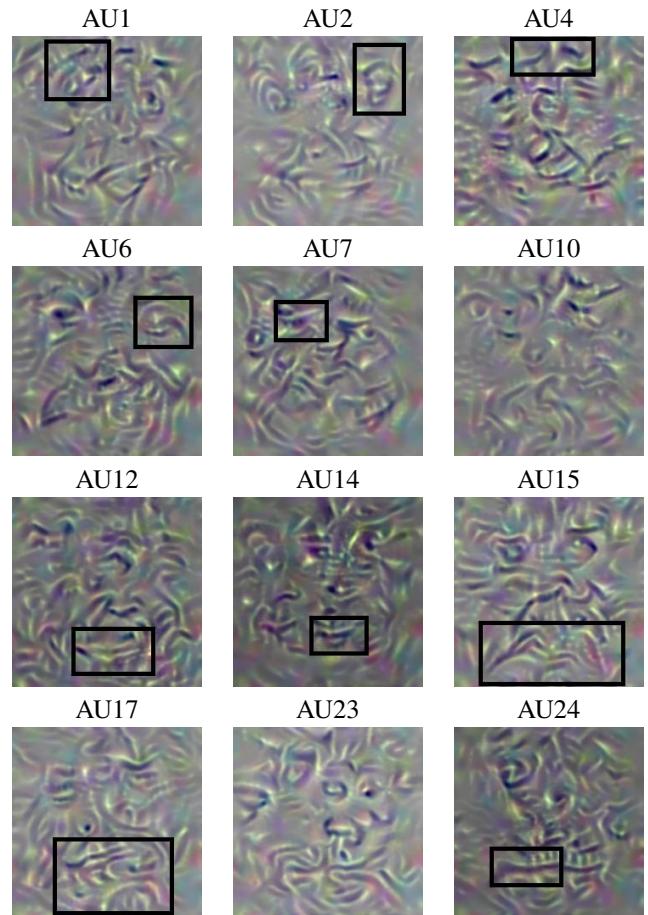
An alternative method for visualizing capsule features is by activation maximization. In this approach the input image is treated as a learnable layer, and changed by gradient ascent for a particular optimization function. This optimization function can be the activation of a neuron in the network, the norm of a layer in the network, or any variations and combinations of the two. The core idea is that as the network modifies the input image to increase the objective function, say the activation of a neuron, the input image will begin to show the visual attribute that the neuron has learned to encapsulate. This process is popularly referred to as ‘deep dreaming’ - see [?] for an excellent overview.

Since the output of a capsule is a vector and not a scalar, we do gradient ascent on its magnitude which is analogous to doing gradient ascent on its activation. A naive applica-



594
595
596
597
598
599
600
601
602
603
604
605
606
607
608
609
610
611
612
613
614
615
616
617
618
619
620
621
622
623
624
625
626
627
628
629
630
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647

Figure 6. Deep dreaming expressions. The top row shows the mean image of each class, and the bottom row shows dreamed images for the class. Line correspond to lines that appear for each expression.



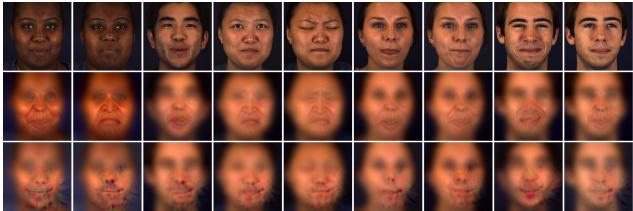
631
632
633
634
635
636
637
638
639
640
641
642
643
644
645
646
647

Figure 7. Deep dreaming AUs. The boxes show where we see the AUs. AU lookup table is Figure ???. For AU 10 and AU 23 the results are difficult to interpret.

tion of this method is prone to create high frequency and non-sensical images. It is therefore necessary to regularize the input image. We use gaussian blurring and random jittering.

4.5.1 Class Capsules

We first visualize the class capsules’ properties for both the expression and action unit detection networks. For both net-



648
649
650
651
652
653
654
655
656
657
658
659
660
661
662
663
664
665
666
667
668
669
670
671
672
673
674
675
676
677
678
679
680
681
682
683
684
685
686
687
688
689
690
691
692
693
694
695
696
697
698
699
700
701

Figure 8. Comparing convolution and capsule units. The top row shows the input images. The second row shows the result of performing gradient ascent visualization on capsule 21 on the mouth area. The bottom row shows the results on the same input images with convolution unit 498 on a finetuned convolution network. While a single capsule can have high magnitude with many different types of mouth positions, a single convolution capsule is only able to maximize activation for a thin upturned corner type mouth, regardless of the type of input image.

works we input a random noise image, and follow the procedure above to amplify the magnitude of each class in turn.

Figure 6 shows the results on the expression network. While the results are not natural - even scary - looking, they demonstrate that the expression capsules have indeed learned to identify correct attributes. For surprise the dark circle below the lips indicates that the capsules have learned to identify a jaw drop. The forehead lines indicate that the network is successfully identifying raised eyebrows by looking for forehead lines. Similarly, for happiness, the diagonal lines moving outward from the mouth corners resemble the shape of smiling lips, and the lines going diagonally down from each side of the nose resemble smile lines.

Figure 7 shows results on a VGG-Face finetuned action unit detection capsule network, with the attributes of the action unit surrounded with a black bounding box. The indicate that for almost every action unit the network is learning the correct characteristics. For example, for AU1, the network not only learns the inside raised eye brow shape, but also the corresponding folds in the skin around it. For AU 24, the network learns that the lips appear narrower and that lines appear around the mouth. For AU 6 the network focuses on the outside of the eyes where laugh lines appear as the outer cheek muscles are drawn up.

At the same time these results are enlightening because they indicate that the network has not completely decorrelated co-occurring but non-causal appearance changes. For example for AU 2, open and smiling mouths are also reconstructed by the network. AU 2 commonly occurs when people are surprised or happy, and so the corresponding mouth changes have also been learned by the network. These results indicate that the network is prone to falsely predict AU 2 when someone is laughing even if their eyebrows are not raised.



702
703
704
705
706
707
708
709
710
711
712
713
714
715
716
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755

Figure 9. Activation maximization visualizing primary capsules. The first and third row show input test images. The second and fourth row show the result of activation maximizing for primary capsule 21 and 5 respectively. The results show that a single capsule is able to model multiple deformations of the mouth area and eyebrows, and activates differently for different action units.



734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755

Figure 10. Activation maximization on the mouth area for capsule 5. While this capsule is able to model changes in brow position well (Figure 9) it is unable to model different mouth positions and is ‘sticky’ towards a downturned mouth position.

4.5.2 Primary Capsules

In Section 4.3 we show that capsule magnitude is linked to how pronounced a expression or AU is, and its direction can be linked to different attributes a face or picture may have. This is done through the reconstruction network. However, class capsules are fully connected to the layer below and therefore do not help us understand how capsules are able to model local part deformations.

To gain insight into how capsules model changes in local parts of the input image, we perform activation maximization on the magnitude of primary capsules’ activations. More specifically, for each action unit, we use gradient ascent to maximize the capsule activation magnitude at a specific activation map location. This process is repeated for all action units. Unlike in Section 4.5.1, we do not use input images that are random noise since that does not let us control the direction in which capsules are activated, and then magnified. Furthermore, since primary capsules are applied to each image in a sliding window fashion, each capsule may have very different responses depending on where its receptive window overlaps with the input image.

756 Instead, we use test images that have perfect predictions,
 757 and the highest confidence (or class activation magnitude)
 758 per test subject and per action unit. This allows us to get
 759 primary capsules that have variation in direction naturally,
 760 while also allowing us ignore shortcomings of the network
 761 (perfect predictions), and test across identities.
 762

763 Spatial locations for each action unit are chosen based on
 764 which part of the face each action unit deforms. In Figure 9
 765 we show examples of different action units becoming more
 766 pronounced with gradient ascent on the *same* primary cap-
 767 sule's magnitude. Furthermore, we find that the cosine dis-
 768 tance between activations after performing gradient ascent.
 769 This shows that the same capsule is able to model different
 770 types of deformations as changes in capsule direction.
 771 However, this is not always true for the primary capsules,
 772 and Figure 10 shows examples of a capsule that is unable to
 773 model deformations around the mouth, but is, on the other
 774 hand, able to model changes around the brow.
 775

776 For comparison, in Figure 8 we perform similar activa-
 777 tion maximization visualization on convolution units. We
 778 finetune a VGG-Face network for action unit detection us-
 779 ing the same training data. We use the activations of the last
 780 convolution layer after ReLU and max-pooling for this pur-
 781 pose since it allows us to have comparable receptive win-
 782 dows size as the primary capsules, and lets us ignore convo-
 783 lution units that are turned off by ReLU and therefore irrel-
 784 evant to the final prediction. We then forward the same se-
 785 lected test images through the network and record the con-
 786 volution units with the highest activation at the specific spa-
 787 tial locations in the activation map for each action unit. We
 788 then perform gradient ascent on the activations of these par-
 789 ticular convolution units at selected spatial locations. Since
 790 we use the post ReLU activations, it is not always possi-
 791 ble to generate visualizations for all input test images, since
 792 the activation can be zeroed out. However, we do find that
 793 unlike capsules, individual convolution units are not able
 794 to model dramatically different deformations, and either do
 795 not activate or exaggerate a fixed type of attribute or at-
 796 tributes exclusively. This is not surprising given that con-
 797 volution units have scalar outputs, while each capsules has
 798 a vector output that is capable of representing more com-
 799 plex information. As a result while a single capsule is able
 800 to have high activation with downturned, puckered, open,
 801 and smiling - among others - mouths, the convolution unit
 802 maximizes activation with thin mouths with a sharp right
 803 upturn regardless of the input image.
 804

4.6. To do and questions

805 Add CVPR 2018 related work and baselines. Add disfa
 806 results and improved Ours-VGG results Add results show-
 807 ing overfitting on Ours vggf. Add cosine similarity numbers
 808 in an attractive way.
 809

Questions: How to best show that after activation max,

810 the capsules become more orthogonal when the resulting
 811 image is different, and less orthogonal when it is the same?
 812 Cosine similarity numbers show this, but i'm not sure how
 813 to add them to the figures. I was thinking of a bar under the
 814 images with the cosine similarity values before and after
 815 max.
 816

817 Should we add ECCV18 numbers? Would it count as
 818 concurrent work because one of the wacv deadlines was be-
 819 fore eccv?
 820

I've removed all the localization routing experiments.
 Do you think that's ok?
 821

5. Conclusions

In this paper, we tested the hypothesis that capsule networks are able to model local part deformations in faces well. We tested this hypothesis by using capsules for action unit detection, and found that capsules are indeed able to model action unit activations well. Our results demonstrated state-of-the-art results on action unit detection on two widely-used datasets. While previous work has shown that capsules are able to model global deformations, we showed that capsules can also capture local deformations. This indicates that capsule networks will also be useful for other tasks where parts of an object need to be modeled well - such as fine-grained classification, or human pose estimation and tracking. At the same time, our results can be useful for the action unit and facial expression understanding community. We propose a novel architecture for action unit detection, and push the state-of-the-art numbers in the area.

In the future, we plan to work on automatic animal facial expression understanding. For a setting such as animal facial expression understanding where data is scarce and difficult to both collect and annotate, it becomes critical to work with models such as capsule networks that are able to extract rich feature representations with fewer overall parameters. In addition, the added ease with which capsule properties can be visualized makes capsule networks an appropriate model working with limited, possibly noisily annotated data.

References

- [1] M. S. Bartlett, G. Littlewort, M. Frank, C. Lainscsek, I. Fasel, and J. Movellan. Recognizing facial expression: machine learning and application to spontaneous behavior. In *Computer Vision and Pattern Recognition, 2005. CVPR 2005. IEEE Computer Society Conference on*, volume 2, pages 568–573. IEEE, 2005. 1, 3
- [2] A. Bulat and G. Tzimiropoulos. How far are we from solving the 2d & 3d face alignment problem? (and a dataset of 230,000 3d facial landmarks). In *International Conference on Computer Vision*, 2017. 4

- 864 [3] W.-S. Chu, F. De la Torre, and J. F. Cohn. Selective transfer
865 machine for personalized facial action unit detection.
866 In *Computer Vision and Pattern Recognition (CVPR), 2013 IEEE Conference on*, pages 3515–3522. IEEE, 2013. 2
- 867 [4] W.-S. Chu, F. De la Torre, and J. F. Cohn. Modeling spatial
868 and temporal cues for multi-label facial action unit detection.
869 *arXiv preprint arXiv:1608.00911*, 2016. 4, 5
- 870 [5] C. A. Corneanu, M. Madadi, and S. Escalera. Deep structure
871 inference network for facial action unit recognition. *ECCV*,
872 2018. 3
- 873 [6] H. Ding, S. K. Zhou, and R. Chellappa. Facenet2expnet:
874 Regularizing a deep face recognition net for expression
875 recognition. *Automatic Face and Gesture Recognition (FG)*,
876 2017. 3, 5
- 877 [7] S. Eleftheriadis, O. Rudovic, and M. Pantic. Multi-
878 conditional latent variable model for joint facial action unit
879 detection. In *Proceedings of the IEEE International Conference
880 on Computer Vision*, pages 3792–3800, 2015. 2
- 881 [8] X. Feng, M. Pietikäinen, and A. Hadid. Facial expression
882 recognition based on local binary patterns. *Pattern Recognition
883 and Image Analysis*, 17(4):592–598, 2007. 3
- 884 [9] S. Ghosh, E. Laksana, S. Scherer, and L.-P. Morency. A
885 multi-label convolutional neural network approach to cross-
886 domain action unit detection. In *Affective Computing and In-
887 telligent Interaction (ACII), 2015 International Conference
888 on*, pages 609–615. IEEE, 2015. 2
- 889 [10] A. Gudi, H. E. Tasli, T. M. den Uyl, and A. Maroulis. Deep
890 learning based faces action unit occurrence and intensity es-
891 timation. In *Automatic Face and Gesture Recognition (FG),
892 2015 11th IEEE International Conference and Workshops on*, volume 6, pages 1–5. IEEE, 2015. 2
- 893 [11] S. Han, Z. Meng, A.-S. Khan, and Y. Tong. Incremental
894 boosting convolutional neural network for facial action unit
895 recognition. In *Advances in Neural Information Processing
896 Systems*, pages 109–117, 2016. 2
- 897 [12] S. Han, Z. Meng, Z. Li, J. O'Reilly, J. Cai, X. Wang, and
898 Y. Tong. Optimizing filter size in convolutional neural net-
899 works for facial action unit recognition. In *The IEEE Confer-
900 ence on Computer Vision and Pattern Recognition (CVPR)*,
901 June 2018. 3
- 902 [13] S. Hinton, Geoffrey E and Sabour and N. Frosst. Matrix
903 capsules with em routing. In *ICLR*, 2018. 1, 2
- 904 [14] S. Jaiswal and M. Valstar. Deep learning the dynamic ap-
905 pearance and shape of facial action units. In *Applications of
906 Computer Vision (WACV), 2016 IEEE Winter Conference on*,
907 pages 1–8. IEEE, 2016. 2, 4
- 908 [15] H. Jung, S. Lee, J. Yim, S. Park, and J. Kim. Joint fine-tuning
909 in deep neural networks for facial expression recognition. In
910 *Proceedings of the IEEE International Conference on Com-
puter Vision*, pages 2983–2991, 2015. 3
- 911 [16] P. Khorrami, T. Paine, and T. Huang. Do deep neural net-
912 works learn facial action units when doing expression recog-
913 nition? In *Proceedings of the IEEE International Conference
914 on Computer Vision Workshops*, pages 19–27, 2015. 3, 5
- 915 [17] Y. Kim, B. Yoo, Y. Kwak, C. Choi, and J. Kim. Deep
916 generative-contrastive networks for facial expression recog-
917 nition. *arXiv preprint arXiv:1703.07140*, 2017. 3, 5
- 918 [18] S. Koelstra, M. Pantic, and I. Patras. A dynamic texture-
919 based approach to recognition of facial actions and their tem-
920 poral models. *IEEE transactions on pattern analysis and
921 machine intelligence*, 32(11):1940–1954, 2010. 2
- 922 [19] W. Li, F. Abtahi, and Z. Zhu. Action unit detection with
923 region adaptation, multi-labeling learning and optimal tem-
924 poral fusing. In *The IEEE Conference on Computer Vision
925 and Pattern Recognition (CVPR)*, July 2017. 1, 2, 4
- 926 [20] Y. Li, J. Chen, Y. Zhao, and Q. Ji. Data-free prior model for
927 facial action unit recognition. *IEEE Transactions on affective
928 computing*, 4(2):127–141, 2013. 2
- 929 [21] M. Liu, S. Li, S. Shan, and X. Chen. Au-aware deep
930 networks for facial expression recognition. In *Automatic
931 Face and Gesture Recognition (FG), 2013 10th IEEE Inter-
932 national Conference and Workshops on*, pages 1–6. IEEE,
933 2013. 3, 5
- 934 [22] M. Liu, S. Li, S. Shan, and X. Chen. Au-inspired deep net-
935 works for facial expression feature learning. *Neurocompu-
936 ting*, 159:126–136, 2015. 3, 5
- 937 [23] P. Liu, S. Han, Z. Meng, and Y. Tong. Facial expression
938 recognition via a boosted deep belief network. In *Pro-
939 ceedings of the IEEE Conference on Computer Vision and Pattern
940 Recognition*, pages 1805–1812, 2014. 3
- 941 [24] P. Lucey, J. F. Cohn, T. Kanade, J. Saragih, Z. Ambadar, and
942 I. Matthews. The extended cohn-kanade dataset (ck+): A
943 complete dataset for action unit and emotion-specified ex-
944 pression. In *Computer Vision and Pattern Recognition Work-
945 shops (CVPRW), 2010 IEEE Computer Society Conference
946 on*, pages 94–101. IEEE, 2010. 5
- 947 [25] S. M. Mavadati, M. H. Mahoor, K. Bartlett, P. Trinh, and J. F.
948 Cohn. Disfa: A spontaneous facial action intensity database.
949 *IEEE Transactions on Affective Computing*, 4(2):151–160,
950 2013. 4
- 951 [26] A. Mollahosseini, D. Chan, and M. H. Mahoor. Going deeper
952 in facial expression recognition using deep neural networks.
953 In *Applications of Computer Vision (WACV), 2016 IEEE
954 Winter Conference on*, pages 1–10. IEEE, 2016. 3
- 955 [27] O. M. Parkhi, A. Vedaldi, A. Zisserman, et al. Deep face
956 recognition. 3
- 957 [28] G. Peng and S. Wang. Weakly supervised facial action
958 unit recognition through adversarial training. In *The IEEE
959 Conference on Computer Vision and Pattern Recognition
960 (CVPR)*, June 2018. 3
- 961 [29] S. Sabour, N. Frosst, and G. E. Hinton. Dynamic routing
962 between capsules. In *Advances in Neural Information Pro-
963 cessing Systems*, pages 3859–3869, 2017. 1, 2, 3
- 964 [30] Z. Shao, Z. Liu, J. Cai, and L. Ma. Deep adaptive atten-
965 tion for joint facial action unit detection and face alignment.
966 *ECCV*, 2018. 3
- 967 [31] K. Sikka, T. Wu, J. Susskind, and M. Bartlett. Exploring bag
968 of words architectures in the facial expression domain. In
969 *Computer Vision–ECCV 2012. Workshops and Demonstra-
970 tions*, pages 250–259. Springer, 2012. 3
- 971 [32] Y. Tong, W. Liao, and Q. Ji. Facial action unit recognition by
972 exploiting their dynamic and semantic relationships. *IEEE
973 transactions on pattern analysis and machine intelligence*,
974 29(10), 2007. 2

- 972 [33] M. F. Valstar, T. Almaev, J. M. Girard, G. McKeown,
973 M. Mehu, L. Yin, M. Pantic, and J. F. Cohn. Fera 2015-
974 second facial expression recognition and analysis challenge.
975 In *Automatic Face and Gesture Recognition (FG), 2015*
976 *11th IEEE International Conference and Workshops on*, volume 6, pages 1–8. IEEE, 2015. 2
- 977 [34] Z. Wang, Y. Li, S. Wang, and Q. Ji. Capturing global semantic
978 relationships for facial action unit recognition. In *Proceedings of the IEEE International Conference on Computer
979 Vision*, pages 3304–3311, 2013. 2
- 980 [35] H. Yang, U. Ciftci, and L. Yin. Facial expression recognition
981 by de-expression residue learning. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June
982 2018. 3
- 983 [36] J. Zeng, W.-S. Chu, F. De la Torre, J. F. Cohn, and Z. Xiong.
984 Confidence preserving machine for facial action unit detection.
985 In *Proceedings of the IEEE International Conference on Computer Vision*, pages 3622–3630, 2015. 2, 4
- 986 [37] X. Zhang and M. H. Mahoor. Simultaneous detection of multiple
987 facial action units via hierarchical task structure learning.
988 In *Pattern Recognition (ICPR), 2014 22nd International Conference on*, pages 1863–1868. IEEE, 2014. 2
- 989 [38] X. Zhang, L. Yin, J. F. Cohn, S. Canavan, M. Reale,
990 A. Horowitz, P. Liu, and J. M. Girard. Bp4d-spontaneous:
991 a high-resolution spontaneous 3d dynamic facial expression
992 database. *Image and Vision Computing*, 32(10):692–706,
993 2014. 4
- 994 [39] Y. Zhang, W. Dong, B.-G. Hu, and Q. Ji. Classifier learning
995 with prior probabilities for facial action unit recognition.
996 In *The IEEE Conference on Computer Vision and Pattern
997 Recognition (CVPR)*, June 2018. 3
- 998 [40] Y. Zhang, W. Dong, B.-G. Hu, and Q. Ji. Weakly-supervised
999 deep convolutional neural network learning for facial action
1000 unit intensity estimation. In *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2018. 3
- 1001 [41] Y. Zhang, R. Zhao, W. Dong, B.-G. Hu, and Q. Ji. Bilateral
1002 ordinal relevance multi-instance regression for facial action
1003 unit intensity estimation. In *CVPR*, June 2018. 3
- 1004 [42] K. Zhao, W.-S. Chu, F. De la Torre, J. F. Cohn, and H. Zhang.
1005 Joint patch and multi-label learning for facial action unit detection.
1006 In *Proceedings of the IEEE Conference on Computer Vision and Pattern
1007 Recognition*, pages 2207–2216, 2015. 2, 4
- 1008 [43] K. Zhao, W.-S. Chu, and H. Zhang. Deep region and multi-label
1009 learning for facial action unit detection. In *Proceedings of the IEEE Conference on Computer Vision and Pattern
1010 Recognition*, pages 3391–3399, 2016. 1, 2, 4
- 1011 [44] X. Zhao, X. Liang, L. Liu, T. Li, Y. Han, N. Vasconcelos,
1012 and S. Yan. Peak-piloted deep network for facial expression
1013 recognition. In *European Conference on Computer Vision*,
1014 pages 425–442. Springer, 2016. 3
- 1015 [45] L. Zhong, Q. Liu, P. Yang, B. Liu, J. Huang, and D. N.
1016 Metaxas. Learning active facial patches for expression analysis.
1017 In *Computer Vision and Pattern Recognition (CVPR), 2012 IEEE Conference on*, pages 2562–2569. IEEE, 2012. 3
- 1018 [46] Y. Zhu, S. Wang, L. Yue, and Q. Ji. Multiple-facial action
1019 unit recognition by shared feature learning and semantic
1020 relation modeling. In *Pattern Recognition (ICPR), 2014 22nd
1021 International Conference on*, pages 1663–1668. IEEE, 2014. 2
- 1022 [47] Y. Zhu, S. Wang, L. Yue, and Q. Ji. Multi-facial action
1023 unit recognition by shared feature learning and semantic
1024 relation modeling. In *Pattern Recognition (ICPR), 2014 22nd
1025 International Conference on*, pages 1663–1668. IEEE, 2014. 2